

Interactive Interview Session

[top product-based companies]

AppliedAIcourse.com

- Sample of Questions & Variations.
- 1 or 2 rounds of programming
- 4 - 5 rounds of DS | ML | DL | Math [depth of understanding]
 - ↳ Depth, breadth, Foundational Maths, Applied, Problem-solving etc.
- Submit your answers / ideas in the chat window

→ role-dependent [DA, DS, MLS, MLE, SDE-ML]

→ follow-up questions

→ behavioral (and) past-experience (and) projects based

→ open-ended questions : build Alexa's wake word detection
[skipped]

→ skipped straight-forward questions (e.g. explain GBDT)

What do companies want to test?

- ① Depth, breadth & clarity of understanding of core concepts in Math, DS, ML, DL
- ② Problem-solving
- ③ Ability to write code
- ④ Tackle open-ended problems [experienced | senior roles]

Q1 . Implement a function to merge K-sorted arrays/lists.
Explain it's time and space complexity

Microsoft, Amazon, Netflix etc.

Pseudo-code/Logic: 5-10 mins
+
Implementation : ~15 mins

L_1

2	3	3	5	6
---	---	---	---	---

L_2


1	7	8	9
---	---	---	---

L_3

0	2	4	6
---	---	---	---

L_4

0	2	2	4	5	6
---	---	---	---	---	---



Time - Complex:

Space - Complex:

Variations:

→ optimized find-min [CS]


→ Merge sorted very large files [CS]

Q2 . Given a sample of observations, how would you estimate the 99th percentile value of the population?

Google Data Science

Central-Limit-Theorem

Non-parametric stats:



Variations :

- ① estimate Median - CLT for percentiles
[hardcore stats]
 - ② estimate uncertainty in model weights
- Bootstrapping Bayesian models.
[BNN](#)

[Source](#)

Q3 . Given the following statistics, what is the probability that a woman has cancer if she has a positive mammogram result?

- One percent of women over 50 have breast cancer.
- Ninety percent of women who have breast cancer test positive on mammograms.
- Eight percent of women will have false positives.

Most Data Science roles across almost all companies

Practice problems:

<https://www.statisticshowto.com/bayes-theorem-problems/>

<http://wwwf.imperial.ac.uk/~atw/Bayes.pdf>

<http://gtribello.github.io/mathNET/bayes-theorem-problems.html>

$$P(C|G) = 0.01$$

$$P(+|C) = 0.90$$

$$P(+|NC) = 0.08$$

calc: $P(+|C) = \frac{P(C \cap +)}{P(+)}$

$$= \frac{P(+|C) P(C)}{P(+ \cap C) P(C \cup + \cap NC)}$$

$$= \frac{0.9 \times 0.01}{(0.9 \times 0.01) + (0.08 \times 0.99)}$$

$$= 0.10 = 10\%$$

Q4. Implement a function `SGDLinearRegression()` using basic NumPy/SciPy/Python which has the following properties:

1. Input : Dataset, lambda(multiplier to the regularization term), batch-size, #epochs
2. Implement Linear regression with Squared loss and L1-regularization
3. Implement batch-SGD optimizer
4. Output: trained Weight-vector

Pseudo-code: 5-10 mins

+

Python implementation: ~20 mins

Amazon, Netflix & many startups

$$\mathcal{L} = \min_{\omega} \sum_{j=1}^n (y_j - \omega^T x_j)^2 + \|\omega\|_1 \rightarrow |w_1| + |w_2| + \dots + |w_d|$$

$$\frac{\partial L}{\partial w_i} = 2 \cdot \sum_{j=1}^n (y_j - \omega^T x_j) \cdot (-x_{ji}) + \underbrace{\text{sign}(w_i)}$$

$$\begin{cases} +1 & w_i > 0 \\ -1 & w_i < 0 \\ 0 & w_i = 0 \end{cases}$$

$$w_i^{\text{new}} = w_i^{\text{old}} - \eta \frac{\partial L}{\partial w_i}$$

$$\eta = 0.1$$

standardize(\mathcal{D})

$w = \text{random-vector}(d)$

for i in $1, 2, \dots, \text{epochs}$

for b in $\text{Batch}(\mathcal{D}, bs)$

for j in $1, 2, \dots, d$

$$w_j := \eta * \frac{\partial L}{\partial w_j}(\mathcal{D}, w, j)$$

Update(η)

Variations:


- logistic regression or L₂-SVM
- elastic-Net reg.
- custom-loss function
- Adam optimizer instead of SGD

Q5. Why is RBF kernel the most popular & default kernel in SVMs, especially when we don't have custom-kernels for the problem at hand

Facebook/Google
ML roles

$$K(x_i, x_j) = \exp \left\{ - \frac{\|x_i - x_j\|^2}{2\sigma^2} \right\}$$


$x_i, x_j \in \mathbb{R}^d$



Variations :

- ① Why do margin & regularization look similar?
- ② Why are SVMs not as popular as GBDT / RF?


Linear SVM:



e.g: Oracle/Yahoo!

Q6. Which of these points have a non-zero slack in the Soft-margin Linear-SVM's primal formulation.

Linear SVM:



$$\min_{w, \xi_i} \frac{\|w\|}{2} + C \sum_{i=1}^n \xi_i$$

$$\text{s.t. } \boxed{1 - y_i(w^T x_i + b)} \geq \xi_i \quad \forall i$$
$$\xi_i \geq 0 \quad \forall i$$

$$\max(0, 1 - y_i(w^T x_i + b))$$

Q7. While constructing GBDT, how do you compute the multiplier term for each base-learner

$$F_m(x) = F_{m-1}(x) + \gamma_m h_m(x).$$

Model using
first 'm' base-learners

↳ m^{th} base-learner

$$F_m(x) = F_{m-1}(x) + \gamma_m h_m(x).$$

already known

Model already
fit on pseudo-residuals

3. Compute multiplier γ_m by solving the following **one-dimensional optimization** problem:

$$\gamma_m = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, F_{m-1}(x_i) + \gamma h_m(x_i)).$$

constant

constant

Variations

① step-by-step GBDT pseudo-code

② modify GBDT to build K base-learners
in each step.

③ Distributed GBDT [senior roles]

Q8. Explain step-by-step (or via pseudo-code), how backpropagation occurs through a batch-normalization layer? [15-min]

e.g: Flipkart, Microsoft

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;
 Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i$ // mini-batch mean

$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2$ // mini-batch variance

$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$ // normalize

$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i)$ // scale and shift

Output

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.

original paper

[Forward-pass]

Backward pass:

simple derivatives
& chain rule

https://kevinzakka.github.io/2016/09/14/batch_normalization/

<https://kratzert.github.io/2016/02/12/understanding-the-gradient-flow-through-the-batch-normalization-layer.html>

Variations

1. explain backprop step by step [30 min]
2. backprop with Dropout | Maxpool / ...
3. L₁ reg of weights in a layer ?

Q9. Explain the training of a GAN, preferably with pseudo-code.

[15 min]

Amazon, Google, MSR

- depends on how you set-up the problem
- multiple correct alternatives

<https://www.tensorflow.org/tutorials/generative/dcgan>

<https://machinelearningmastery.com/how-to-code-the-generative-adversarial-network-training-algorithm-and-loss-functions/>


Variations

1. LSTM training

2. Transformer <https://www.tensorflow.org/tutorials/text/transformer>

Q10. You have an ML classification model, M, running in production trained on data from days 1 to 10. Today is day 25. What metric(s) would you keep track of, to update the production model using more recent data.

Amazon, Microsoft, Facebook



→ Runtime dataset difference across days

The hardest questions in most interviews are based around “Why?”

56 Sample Interview/Conceptual Questions with Audio explanation:

<https://www.appliedaicourse.com/lecture/11/applied-machine-learning-online-course/4209/sample-interview-and-conceptual-questions-audio/69/module-10-live-sessions-on-miscellaneous-topics-audio-answers>